

COMSM0045: PRACTICAL-Lab4

Dima Damen

`Dima.Damen@bristol.ac.uk`

Bristol University, Department of Computer Science
Bristol BS8 1UB, UK

November 10, 2021

What else can we do?

- ▶ Data Augmentation
- ▶ Debugging Strategies
- ▶ Dropout

Data Augmentation

- ▶ Data augmentation is making the most of the training samples by introducing variations of these samples to accommodate for required invariances

Data Augmentation

- ▶ Data augmentation is making the most of the training samples by introducing variations of these samples to accommodate for required invariances
- ▶ Why Data Augmentation?

Data Augmentation

- ▶ Data augmentation is making the most of the training samples by introducing variations of these samples to accommodate for required invariances
- ▶ Why Data Augmentation?
 - ▶ Because it's all about the size of your data → More data for training

Data Augmentation

- ▶ Data augmentation is making the most of the training samples by introducing variations of these samples to accommodate for required invariances
- ▶ Why Data Augmentation?
 - ▶ Because it's all about the size of your data → More data for training
 - ▶ **More importantly...** to accommodate invariances

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint
Number plate recognition	

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint
Number plate recognition	translation, scaling

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint
Number plate recognition	translation, scaling
Action Recognition	

Invariances in data

- ▶ A *problem* is **invariant** to a *property* when the problem remains unchanged when transformations of a certain type are applied.

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint
Number plate recognition	translation, scaling
Action Recognition	translation, rotation, scaling, viewpoint, speed

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- How to provide invariance?

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- ▶ How to provide invariance? → *artificially* augment for:
 - ▶ Translation:
 - ▶ Rotation:
 - ▶ Scaling:
 - ▶ Viewpoint:

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- ▶ How to provide invariance? → *artificially* augment for:
 - ▶ Translation: shifts – luckily CNNs do that for us
 - ▶ Rotation:
 - ▶ Scaling:
 - ▶ Viewpoint:

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- ▶ How to provide invariance? → *artificially* augment for:
 - ▶ Translation: shifts – luckily CNNs do that for us
 - ▶ Rotation: rotations
 - ▶ Scaling:
 - ▶ Viewpoint:

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- ▶ How to provide invariance? → *artificially* augment for:
 - ▶ Translation: shifts – luckily CNNs do that for us
 - ▶ Rotation: rotations
 - ▶ Scaling: croppings
 - ▶ Viewpoint:

Invariances in data

Problem	Invariant to...
Object Recognition	translation, rotation, scaling, viewpoint

- ▶ How to provide invariance? → *artificially* augment for:
 - ▶ Translation: shifts – luckily CNNs do that for us
 - ▶ Rotation: rotations
 - ▶ Scaling: croppings
 - ▶ Viewpoint: Minor - affine transformations, otherwise :-(collect more data!

Invariances in data

- ▶ Other invariances:
 - ▶ invariance to random noise
 - ▶ invariance to occlusion
 - ▶ invariance to lighting conditions
 - ▶ invariance to colour variations
 - ▶ invariance to time of year!? — Generative!

Data Augmentation

- Data augmentation for invariances existed before deep learning




Fig. 7. Warped patches from the images of Fig. 6 show the range of

Data Augmentation

- ▶ Why can't current deep learning methods do that automatically for us?

Data Augmentation

- Crazier ideas are starting to show up in data augmentation... For example, methods newly use Mixup



Cat

* 0.5 +



Dog

* 0.5 =



50%: Cat, 50%: Dog ,

¹ Fig from <https://amitness.com/2020/05/data-augmentation-for-nlp/>

Data Augmentation

- Or CutMix

CutMix



Dog 0.6

Cat 0.4

2

²Fig from Yun et al (2019). CutMix: Regularization Strategy to Train Strong Classifiers with Localizable Features

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly
 - ▶ quite well

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly
 - ▶ quite well

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well → you're done! Find a more interesting problem

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well → you're done! Find a more interesting problem
- ▶ I collected more data but things did not improve?

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well → you're done! Find a more interesting problem
- ▶ I collected more data but things did not improve?
 - ▶ Think about the *quality* of your data

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well → you're done! Find a more interesting problem
- ▶ I collected more data but things did not improve?
 - ▶ Think about the *quality* of your data
 - ▶ Think about the *quality* of your *labels*

But... Do I need more Data?

- ▶ There is a balance between the expense of collecting labelled data and refining the method
- ▶ How are you performing on your training data?
 - ▶ poorly → your algorithm needs work, you are not making the most of the data you have
 - ▶ quite well → How are you performing on your test data?
 - ▶ poorly → try augmentation, otherwise collect more data
 - ▶ quite well → you're done! Find a more interesting problem
- ▶ I collected more data but things did not improve?
 - ▶ Think about the *quality* of your data
 - ▶ Think about the *quality* of your *labels*
 - ▶ → fix and start over

Debugging Deep Learning Algorithms

- ▶ When a general machine learning code performs poorly, including deep learning code, it is very tricky to decide whether that is a bug in the code or a problem in the algorithm
- ▶ Compiling correctly and getting numbers out is not an indication of correctness

Debugging Deep Learning Algorithms

- ▶ When a general machine learning code performs poorly, including deep learning code, it is very tricky to decide whether that is a bug in the code or a problem in the algorithm
- ▶ Compiling correctly and getting numbers out is not an indication of correctness
- ▶ We do not know what the "correct" implementation will give in terms of accuracy, that is in fact what we wish to discover
- ▶ Careful debugging is thus a must

What is your baseline performance?

- ▶ Remember: you cannot perform worse than the baseline!

What is your baseline performance?

- ▶ Remember: you cannot perform worse than the baseline!
- ▶ What can an algorithm that makes decisions by “chance” do?

What is your baseline performance?

- ▶ Remember: you cannot perform worse than the baseline!
- ▶ What can an algorithm that makes decisions by “chance” do?
- ▶ For a binary classifier, your baseline is 50%

What is your baseline performance?

- ▶ Remember: you cannot perform worse than the baseline!
- ▶ What can an algorithm that makes decisions by “chance” do?
- ▶ For a binary classifier, your baseline is 50%
- ▶ For a classifier into N balanced classes, your baseline is $\frac{1}{N}\%$

What is your baseline performance?

- ▶ Remember: you cannot perform worse than the baseline!
- ▶ What can an algorithm that makes decisions by “chance” do?
- ▶ For a binary classifier, your baseline is 50%
- ▶ For a classifier into N balanced classes, your baseline is $\frac{1}{N}\%$
- ▶ For a classifier with unbalanced classes??

Debugging Strategies

- ▶ *Mickey Mouse Examples* - test your solution on small tests that you know the outcome for

Debugging Strategies

- ▶ *Mickey Mouse Examples* - test your solution on small tests that you know the outcome for
- ▶ Evaluate the performance of your building blocks in isolation

Debugging Strategies

- ▶ *Mickey Mouse Examples* - test your solution on small tests that you know the outcome for
- ▶ Evaluate the performance of your building blocks in isolation
- ▶ Monitor the model in action

Debugging Strategies

- ▶ *Mickey Mouse Examples* - test your solution on small tests that you know the outcome for
- ▶ Evaluate the performance of your building blocks in isolation
- ▶ Monitor the model in action
- ▶ Look at failure cases (qualitative assessment)

Debugging Strategies

- ▶ *Mickey Mouse Examples* - test your solution on small tests that you know the outcome for
- ▶ Evaluate the performance of your building blocks in isolation
- ▶ Monitor the model in action
- ▶ Look at failure cases (qualitative assessment)
- ▶ Checkpoints and model saving

Debugging Strategies

“Directly observing the machine learning model performing its task will help to determine whether the quantitative performance numbers it achieves seem reasonable”

Goodfellow et al, p432

Dima Damen

Dima.Damen@bristol.ac.uk

COMSM0045: PRACTICAL - Lab4 - 2021/2022

Debugging Strategies

“Evaluation bugs can be some of the most devastating bugs because they can mislead you into believing your system is performing well when it is not”

Goodfellow et al, p432

Dima Damen

Dima.Damen@bristol.ac.uk

COMSM0045: PRACTICAL - Lab4 - 2021/2022

Dropout

- ▶ What is regularisation?
- ▶ Remind yourself about dropout, as a regularisation/ensemble approach, from the lectures.

By the end of these lab sessions, you should be able to...

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture
- ▶ Define a shallow Convolutional Neural Network (CNN) architecture

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture
- ▶ Define a shallow Convolutional Neural Network (CNN) architecture
- ▶ Train and validate a CNN, and monitor its progress and results using Tensorboard

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture
- ▶ Define a shallow Convolutional Neural Network (CNN) architecture
- ▶ Train and validate a CNN, and monitor its progress and results using Tensorboard
- ▶ Understand and estimate the effect of changing hyper-parameters on your results

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture
- ▶ Define a shallow Convolutional Neural Network (CNN) architecture
- ▶ Train and validate a CNN, and monitor its progress and results using Tensorboard
- ▶ Understand and estimate the effect of changing hyper-parameters on your results
- ▶ Implement and evaluate a variety of data augmentation techniques

By the end of these lab sessions, you should be able to...

- ▶ Define a Fully-Connected Deep Neural Network (DNN) architecture
- ▶ Define a shallow Convolutional Neural Network (CNN) architecture
- ▶ Train and validate a CNN, and monitor its progress and results using Tensorboard
- ▶ Understand and estimate the effect of changing hyper-parameters on your results
- ▶ Implement and evaluate a variety of data augmentation techniques
- ▶ Implement dropout as one of the most common regularisation approaches

And now....

READY....

STEADY....

GO...